

# Computational Photography

Matthias Zwicker, Siavash Bigdeli  
University of Bern  
Fall 2014






# Project 4

- Poisson image editing
- Image segmentation using graph cut optimization

# Poisson image editing



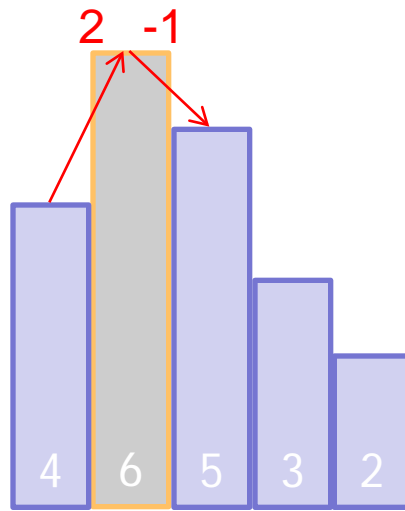

- Implement a Poisson solver
- Demonstrate using different scenarios
  - Seamless cloning
  - Combination of images
  - Highlight removal (gamma compression)

# Poisson Image Editing




?

# Poisson Image Editing



After many iterations



# Reminder


- Goal: find image  $f$  such that

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}$$

- **Linear least squares** problem


$$\min_{f|_{\Omega}} \sum_{\langle p, q \rangle \cap \Omega \neq \emptyset} \underbrace{(f_p - f_q)}_{\text{x or y component of discrete gradient}} - \underbrace{v_{pq}}_{\text{x or y component of guidance fields}})^2, \text{ with } f_p = f_p^*, \text{ for all } p \in \partial\Omega$$



x or y component  
of discrete gradient



Pixel neighborhood

x or y component  
of guidance fields




  $f^*$  known  
  $f$  unknown

# Gauss-Seidel solver

- Loop many times
  - Loop over all unknown pixels  $p$ 
    - At each pixel, solve for pixel value  $f_p$  using **current values** for neighboring pixels  $f_q$


$$|N_p| f_p - \sum_{q \in N_p \cap \Omega} f_q = \sum_{q \in N_p \cap \partial\Omega} f_q^* + \sum_{q \in N_p} v_{pq}$$



- Pixel neighborhood



$$q \in N_p$$

$$|N_p| = 4$$



  $f^*$  known  
  $f$  unknown

# Gauss-Seidel solver

- Loop many times
  - Loop over all unknown pixels  $p$ 
    - At each pixel, solve for pixel value  $f_p$  using **current values** for neighboring pixels  $f_q$

$$|N_p|f_p - \sum_{q \in N_p \cap \Omega} f_q = \sum_{q \in N_p \cap \partial\Omega} f_q^* + \sum_{q \in N_p} v_{pq}$$

- Thousand(s) of iterations necessary for convergence on large images  
-> very slow
- **Test with small images!**



# Speed up Gauss-Seidel solver

- Work only on affected pixels

```
% Find affected pixels
```

```
[I,J] = find(mask(:, :, 1) == 0);
```

```
for it=1:nIter
```

```
    % Work only on found pixels
```

```
    for k=1:length(I)
```

```
        i = I(k);
```

```
        j = J(k);
```

```
        out(i, j) = ...
```

```
    end
```

```
end
```

# Speed up Gauss-Seidel solver

- Downsampling


```
% Half of the steps "down sampled"
nIter = ceil(nIter/2);
% downsampling for the first half
outDS = imresize(out, 0.5);
maskDS = imresize(mask, 0.5);
[vxDs vyDs] = imageGradient(imresize(source, 0.5));
% Do Gauss-Seidel iterations
for k=1:nIter
    ...
end

% upsampling for the second half.
out = imresize(outDS, 2);
% Do Gauss-Seidel iterations
for k=1:nIter
    ...
end
```

# Implementation detail

- $v_{pq}$  includes sign according to order of  $pq$ !
- Desired gradient field stored in  $v_x, v_y$

$$\sum_{q \in N_p} v_{pq} = v_y(i-1, j) - v_y(i, j) + v_x(i, j-1) - v_x(i, j)$$



- Test implementation by reconstructing image without modifying gradient

# Gradient

- Matlab provides a function to calculate gradients: `Gradient(A)`. **Dont use it!**
- Use instead  
`imfilter(A, [-1 1], 'same')`
- Or  
`A(:, 2:end) - A(:, 1:end-1)`

# Gradient

- `Gradient(A)` vs. `A(:,2:end)-A(:,1:end-1)`

A =

>> A(2:10)-A(1:9)

ans =


A	96	15	97	95	48	80	14	42	91	79
A(2:10)		96	15	97	95	48	80	14	42	91
A(1:9)	96	15	97	95	48	80	14	42	91	
ans		-81	82	-2	-47	32	-66	28	49	-12

$$f'(p) = f(p+1) - f(p)$$

# Gradient

- `Gradient(A)` vs. `A(:,2:end)-A(:,1:end-1)`

A =



$$f'(p) = \frac{1}{2}([f(p) - f(p-1)] + [f(p+1) - f(p)]) = \frac{1}{2}(f(p+1) - f(p-1))$$

# Gradient (A)



$A(:, 2:\text{end}) - A(:, 1:\text{end}-1)$






# Seamless cloning




Target



Source



Mask




Output

# Seamless cloning






# Seamless cloning







# Input to Poisson solver



Target (boundary cond.)




Mask





$x$  and  $y$  components of desired gradient taken from source

# Gradient mixing



- Desired gradient is constructed mixing gradients of two images



Input images




Gradient  $x$  comp.





Gradient  $y$  comp.

# Gradient mixing using maximum

- At each pixel, desired  $x, y$  gradient components contain larger of two values from input images





Input  $y$  comp.



Desired  $y$  comp.



Input  $x$  comp.



Desired  $x$  comp.




# Input to poisson solver




Target



Desired  $y$  comp.




Mask with 1-pixel boundary



Desired  $x$  comp.

# Result



Input images



Result



# Gradient compression

- Desired gradient by gamma compression
- Use gradient  $\nabla f^*$  of logarithm of input image
  - Weber's law: brightness perception is logarithmic  
[http://en.wikipedia.org/wiki/Weber%E2%80%93Fechner\\_law](http://en.wikipedia.org/wiki/Weber%E2%80%93Fechner_law)

- Desired gradient  $\mathbf{v}$


$$\mathbf{v} = \alpha^\beta |\nabla f^*|^{-\beta} \nabla f^*$$

- $\alpha$ : small constant ( $\sim 0.005$ ) times average gradient magnitude
- $\beta$ :  $\sim 0.4$

# Result



Target



Mask



Output

# Project 4

- Poisson image editing
- Image segmentation using graph cut optimization

# Interactive image segmentation

- Four steps
  1. Interactive foreground and background labeling
  2. Construct color model
  3. Set up data and smoothness terms, graph connectivity
  4. Solve optimization using graph cuts

# Interactive labeling

- Use `imfreehand` to draw foreground and background masks on images
- Get interactive keyboard input using `input` to select foreground or background
- Store in binary images
- Visualize the results

# Interactive labeling



User provided masks



Binary representation

# Color model

- Extract labeled foreground and background pixels
- Input image `img`, binary masks `fmask`, `bmask`

```
% get all pixel colors in a 3 x (#pixels) array
colors = reshape(shiftdim(img,2),3,size(img,1)*size(img,2));

% get all pixel coordinates in a 2 x (#pixels) array
x = zeros(size(img,1),size(img,2),2);
[x(:,:,1) x(:,:,2)] = meshgrid(1:size(img,2),1:size(img,1));
x = reshape(shiftdim(x,2),2,size(img,1)*size(img,2));

% get the foreground and background masks in a 1 x (#pixels) array
fmask = reshape(fmask,1,size(img,1)*size(img,2));
bmask = reshape(bmask,1,size(img,1)*size(img,2));

% get the indices of the foreground and background pixels
[tt ttt findices] = intersect(x', (x.*[fmask; fmask])', 'rows');
[tt ttt bindices] = intersect(x', (x.*[bmask; bmask])', 'rows');

% extract only labeled foreground and background pixels
fcolors = colors(:,findices);
bcolors = colors(:,bindices);
```

# Color model

- Use `gmdistribution` to obtain Gaussian mixture model
  - Two components should be sufficient
- Obtain the estimated probability distribution using `pdf`
  - Type `help gmdistribution/pdf`



# Color model




Background probabilities



Foreground probabilities



Means of bgr. components




Means of fgr. components

# Graph cut

- Obtain graph cut code from  
<http://vision.ucla.edu/~brian/gcmex.html>
- Solve graph cut using  
`GCMex(class, unary, pairwise, labelcost, 0)`
- Study example for how to use it
- Assume: image with  $W \times H = N$  pixels


# Parameters

- `class`
  - Initial labeling (red/blue in the graph)
  - $1 \times N$  matrix ( $N = W \times H$ )
  - All zeros
- `unary`
  - Correlation between pixel and labels (data term, black lines in the graph)
  - With 2 labels,  $2 \times N$  matrix
  - Fill in with data gained from color model and user provided masks



# Parameters

- pairwise
  - Smoothness term (part of orange lines)
  - Stored in sparse  $N \times N$  matrix
  - Max. 8 non-zero entries per row
    - Because each pixel has at most 8 neighbors
  - Construct sparse matrix using `sparse`
    - Directly writing into matrix is extremely slow
    - Construct index arrays first, then pass to `sparse(rowindices, colindices, values, N, N);`




# Pairwise

- Each matrix entry is  $e^{-\beta \|c_p - c_q\|^2}$
- Neighboring pixels  $p, q$ , pixel colors  $c_p, c_q$
- Visualization by summing over columns




# Parameters

- labelcost
  - 2×2 matrix for binary labels
  - Values depend of labels of interacting pixels
  - Values are multiplied with matrix entries in pairwise to get smoothness term
  - Use  $\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$



# Result




Background



Foreground



Binary  
labels



Input  
masks


# Bonus

- Interactive digital photomontage
- Only 2 photos
- Combination of the pervious assignments
- New data and smoothness terms for the graph cut
- See project webpage for the original paper and a C++ implementation


<http://grail.cs.washington.edu/projects/photomontage/>



# Bonus




# Bonus





# Bonus



# Bonus

